

# Models for Polytomous Responses

Shreejit

6/5/2020

```
#install.packages("car")  
#install.packages("VGAM")
```

For a target with Ordinal Values

## Proportional Odds Model

Unless the response variable has numeric values, it is important to ensure that it has been defined as an ordered factor (using `ordered()`). In the Arthritis data, the response, Improved was setup this way, as we can check by printing some of the values

```
library(vcd)
```

```
## Loading required package: grid
```

```
library(grid)  
data("Arthritis", package="vcd")  
head(Arthritis$Improved, 8)
```

```
## [1] Some    None    None    Marked Marked Marked None    Marked  
## Levels: None < Some < Marked
```

```
library(MASS)  
arth.polr <- polr(Improved ~ Sex + Treatment + Age, data=Arthritis, Hess=TRUE)  
summary(arth.polr)
```

```
## Call:
## polr(formula = Improved ~ Sex + Treatment + Age, data = Arthritis,
##       Hess = TRUE)
##
## Coefficients:
##               Value Std. Error t value
## SexMale        -1.25168    0.54636  -2.291
## TreatmentTreated  1.74529    0.47589   3.667
## Age              0.03816    0.01842   2.072
##
## Intercepts:
##               Value Std. Error t value
## None|Some      2.5319   1.0571    2.3952
## Some|Marked    3.4309   1.0912    3.1443
##
## Residual Deviance: 145.4579
## AIC: 155.4579
```

Summary gives us the coefficients, Intercepts labeled by the cutpoint on the ordinal response but no p-values. The `car::Anova()` method gives the appropriate tests

```
library(car)
```

```
## Loading required package: carData
```

```
Anova(arth.polr)
```

```
## Analysis of Deviance Table (Type II tests)
##
## Response: Improved
##               LR Chisq Df Pr(>Chisq)
## Sex           5.6880   1  0.0170812 *
## Treatment    14.7095   1  0.0001254 ***
## Age           4.5715   1  0.0325081 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

In R, the PO and NPO models can be readily contrasted by fitting them both using `vglm()` in the VGAM package. This defines the cumulative family of models and allows a parallel option. With `parallel=TRUE`, this is equivalent to the `polr()` model, except that the signs of the coefficients are reversed

```
#install.packages("VGAM")
library(VGAM)
```

```
## Loading required package: stats4
```

```
## Loading required package: splines
```

```
##
## Attaching package: 'VGAM'
```

```
## The following object is masked from 'package:car':
##
##      logit
```

```
arth.po <- vglm(Improved ~ Sex + Treatment + Age, data=Arthritis,family = cumulative(
parallel=TRUE))
arth.po
```

```
##
## Call:
## vglm(formula = Improved ~ Sex + Treatment + Age, family = cumulative(parallel = TR
UE),
##      data = Arthritis)
##
##
## Coefficients:
##      (Intercept):1      (Intercept):2      SexMale TreatmentTreated
##      2.53199006      3.43098785      1.25167130      -1.74530408
##      Age
##      -0.03816279
##
## Degrees of Freedom: 168 Total; 163 Residual
## Residual deviance: 145.4579
## Log-likelihood: -72.72896
```

The more general NPO model can be fit using parallel=FALSE.

```
arth.npo <- vglm(Improved ~ Sex + Treatment + Age, data=Arthritis, family = cumulativ
e(parallel=FALSE))
arth.npo
```

```
##
## Call:
## vglm(formula = Improved ~ Sex + Treatment + Age, family = cumulative(parallel = FA
LSE),
##      data = Arthritis)
##
##
## Coefficients:
##      (Intercept):1      (Intercept):2      SexMale:1
##      2.61853875      3.43117450      1.50982704
##      SexMale:2 TreatmentTreated:1 TreatmentTreated:2
##      0.86643395      -1.83692928      -1.70401146
##      Age:1      Age:2
##      -0.04086648      -0.03729421
##
## Degrees of Freedom: 168 Total; 160 Residual
## Residual deviance: 143.5741
## Log-likelihood: -71.78703
```

The VGAM package defines a `coef()` method that can print the coefficients in a more readable matrix form giving the category cutpoints

```
coef(arth.po, matrix=TRUE)
```

```
##              logitlink(P[Y<=1]) logitlink(P[Y<=2])
## (Intercept)      2.53199006      3.43098785
## SexMale          1.25167130      1.25167130
## TreatmentTreated -1.74530408     -1.74530408
## Age              -0.03816279     -0.03816279
```

```
coef(arth.npo, matrix=TRUE)
```

```
##              logitlink(P[Y<=1]) logitlink(P[Y<=2])
## (Intercept)      2.61853875      3.43117450
## SexMale          1.50982704      0.86643395
## TreatmentTreated -1.83692928     -1.70401146
## Age              -0.04086648     -0.03729421
```

In most cases, nested models can be tested using an `anova()` method, but the VGAM package has not implemented this for “vglm” objects. Instead, it provides an analogous function, `lrtest()`

```
VGAM::lrtest(arth.npo, arth.po)
```

```
## Likelihood ratio test
##
## Model 1: Improved ~ Sex + Treatment + Age
## Model 2: Improved ~ Sex + Treatment + Age
##   #Df  LogLik Df  Chisq Pr(>Chisq)
## 1 160 -71.787
## 2 163 -72.729   3  1.8838    0.5969
```

The LR test can be also calculated as “manually” shown below using the difference in residual deviance for the two models

```
tab <- cbind(Deviance = c(deviance(arth.npo), deviance(arth.po)), df = c(df.residual(a
rth.npo), df.residual(arth.po)))
tab <- rbind(tab, diff(tab))
rownames(tab) <- c("GenLogit", "PropOdds", "LR test")
tab <- cbind(tab, pvalue=1-pchisq(tab[,1], tab[,2]))
tab
```

```
##           Deviance  df    pvalue
## GenLogit 143.574067 160 0.8196626
## PropOdds 145.457915 163 0.8343527
## LR test   1.883849   3 0.5968607
```

The `vglm()` can also fit partial proportional odds models, by specifying a formula giving the terms for which the PO assumption should be taken as TRUE or FALSE. Here we illustrate this using `parallel=FALSE ~ Sex`, to fit separate slopes for males and females, but parallel lines for the other predictors. The same model would be fit using `parallel=TRUE ~ Treatment + Age`.

```
arth.ppo <- vglm(Improved ~ Sex + Treatment + Age, data=Arthritis, family = cumulative
(parallel=FALSE ~ Sex))
coef(arth.ppo, matrix=TRUE)
```

```
##           logitlink(P[Y<=1]) logitlink(P[Y<=2])
## (Intercept)           2.54245216           3.61556073
## SexMale           1.48333554           0.86736192
## TreatmentTreated    -1.77574178          -1.77574178
## Age              -0.03962233          -0.03962233
```

## Graphical Assessment of Proportional Odds

```
#install.packages("polspline",dependencies = TRUE)
#install.packages("rms",dependencies = TRUE)
library("rms")
```

```
## Loading required package: Hmisc
```

```
## Loading required package: lattice
```

```
## Loading required package: survival
```

```
## Loading required package: Formula
```

```
## Loading required package: ggplot2
```

```
##
## Attaching package: 'Hmisc'
```

```
## The following objects are masked from 'package:base':
##
##      format.pval, units
```

```
## Loading required package: SparseM
```

```
##
## Attaching package: 'SparseM'
```

```
## The following object is masked from 'package:base':
##
##      backsolve
```

```
##
## Attaching package: 'rms'
```

```
## The following objects are masked from 'package:VGAM':
##
##      calibrate, lrtest
```

```
## The following objects are masked from 'package:car':
##
## Predict, vif
```

```
arth.po2 <- lrm(Improved ~ Sex + Treatment + Age, data=Arthritis)
arth.po2
```

```
## Logistic Regression Model
```

```
##
## lrm(formula = Improved ~ Sex + Treatment + Age, data = Arthritis)
##
##
```


		Model Likelihood		Discrimination		Rank Discrim.	
		Ratio Test		Indexes		Indexes	
## Obs	84	LR chi2	24.46	R2	0.291	C	0.750
## None	42	d.f.	3	g	1.335	Dxy	0.500
## Some	14	Pr(> chi2)	<0.0001	gr	3.801	gamma	0.503
## Marked	28			gp	0.280	tau-a	0.309
## max  deriv	1e-07			Brier	0.187		

```
##
##
```

	Coef	S.E.	Wald Z	Pr(> Z )
## y>=Some	-2.5320	1.0570	-2.40	0.0166
## y>=Marked	-3.4310	1.0911	-3.14	0.0017
## Sex=Male	-1.2517	0.5464	-2.29	0.0220
## Treatment=Treated	1.7453	0.4759	3.67	0.0002
## Age	0.0382	0.0184	2.07	0.0382

```
##
```

```
op <- par(mfrow=c(1,3))
plot.xmean.ordinaly(Improved ~ Sex + Treatment + Age, data=Arthritis,lwd=2, pch=16, s
ubn=FALSE)
```



```
par(op)
```

## Visualizing results for Proportional Odds Model

```
arth.fitp <- cbind(Arthritis,predict(arth.polr, type="probs"))
head(arth.fitp)
```


##	ID	Treatment	Sex	Age	Improved	None	Some	Marked
## 1	57	Treated	Male	27	Some	0.7326185	0.1380586	0.1293229
## 2	46	Treated	Male	29	None	0.7174048	0.1444325	0.1381627
## 3	77	Treated	Male	30	None	0.7096042	0.1476259	0.1427699
## 4	17	Treated	Male	32	Marked	0.6936286	0.1540035	0.1523679
## 5	36	Treated	Male	46	Marked	0.5702499	0.1950359	0.2347142
## 6	23	Treated	Male	58	Marked	0.4563432	0.2171302	0.3265266



```
library(reshape2)
plotdat <- melt(arth.fitp, id.vars = c("Sex", "Treatment", "Age", "Improved"), measure.
vars=c("None", "Some", "Marked"),
          variable.name = "Level",
          value.name = "Probability")
head(plotdat)
```

```
##      Sex Treatment Age Improved Level Probability
## 1 Male   Treated  27      Some  None    0.7326185
## 2 Male   Treated  29      None  None    0.7174048
## 3 Male   Treated  30      None  None    0.7096042
## 4 Male   Treated  32     Marked  None    0.6936286
## 5 Male   Treated  46     Marked  None    0.5702499
## 6 Male   Treated  58     Marked  None    0.4563432
```

```
library(ggplot2)
library(directlabels)
gg <- ggplot(plotdat, aes(x = Age, y = Probability, colour = Level)) +
  geom_line(size=2.5) + theme_bw() + xlim(10,80) +
  geom_point(color="black", size=1.5) +
  facet_grid(Sex ~ Treatment) #, labeller = function(x, y) sprintf("%s = %s", x, y)
direct.label(gg)
```



Although we now have three response curves in each panel, this plot is relatively easy to understand: (a) In each panel, the probability of no improvement decreases with age, while that for marked improvement increases. (b) It is easy to compare the placebo and treated groups in each row, showing that no improvement decreases, while marked improvement increases with the active treatment. (On the other hand, this layout makes it harder to compare panels vertically for males and females in each condition.) (c) The points show where the observations are located in each panel; so, we can see that the data is quite thin for males given the placebo

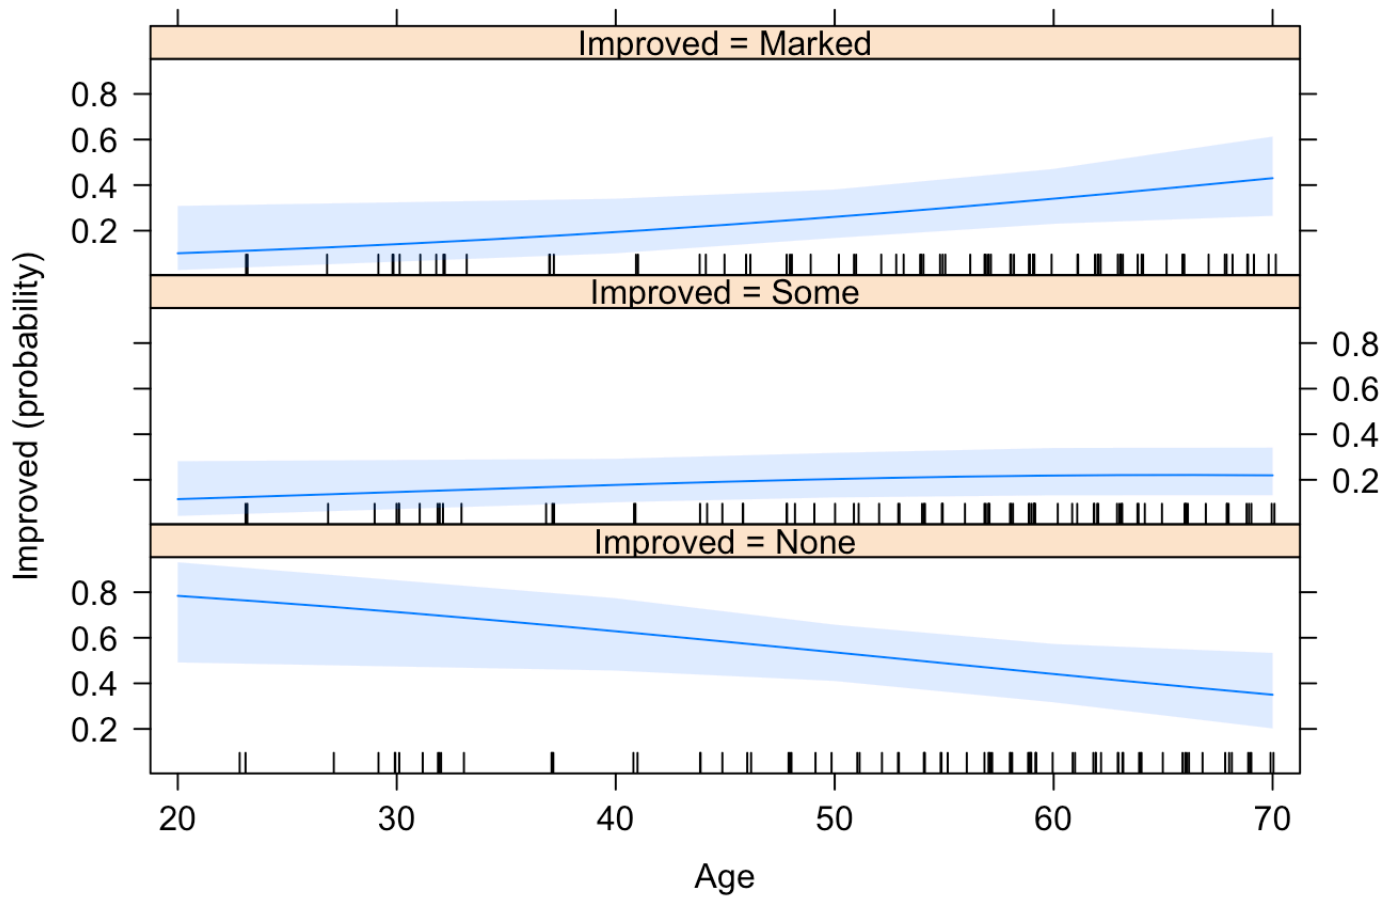
## Effect Plots

```
#install.packages("effects")
library(effects)
```

```
## Use the command
##   lattice::trellis.par.set(effectsTheme())
##   to customize lattice options for effects plots.
## See ?effectsTheme for details.
```

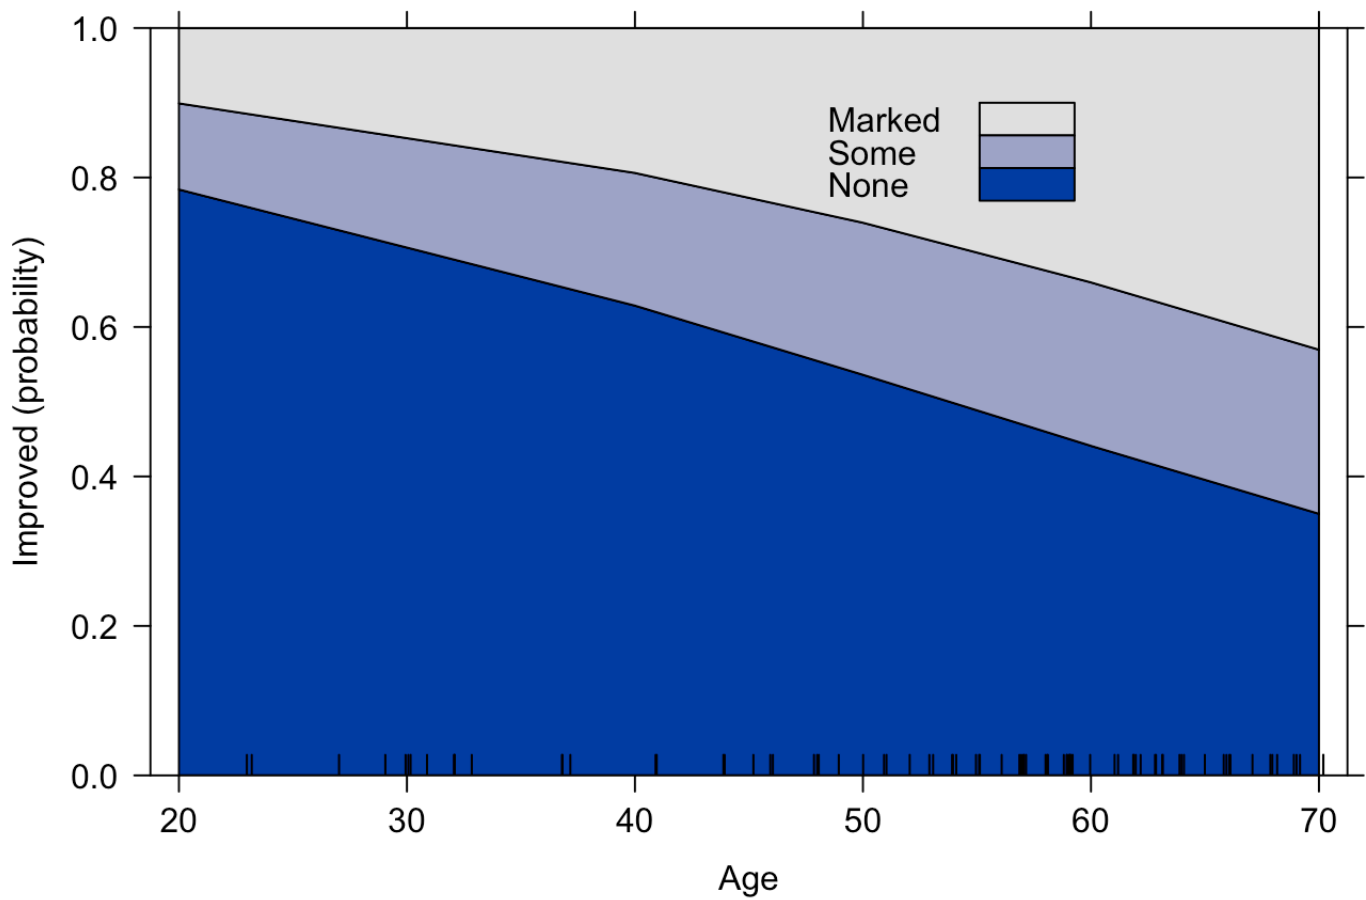
```
plot(Effect("Age", arth.polr))
```

### Age effect plot

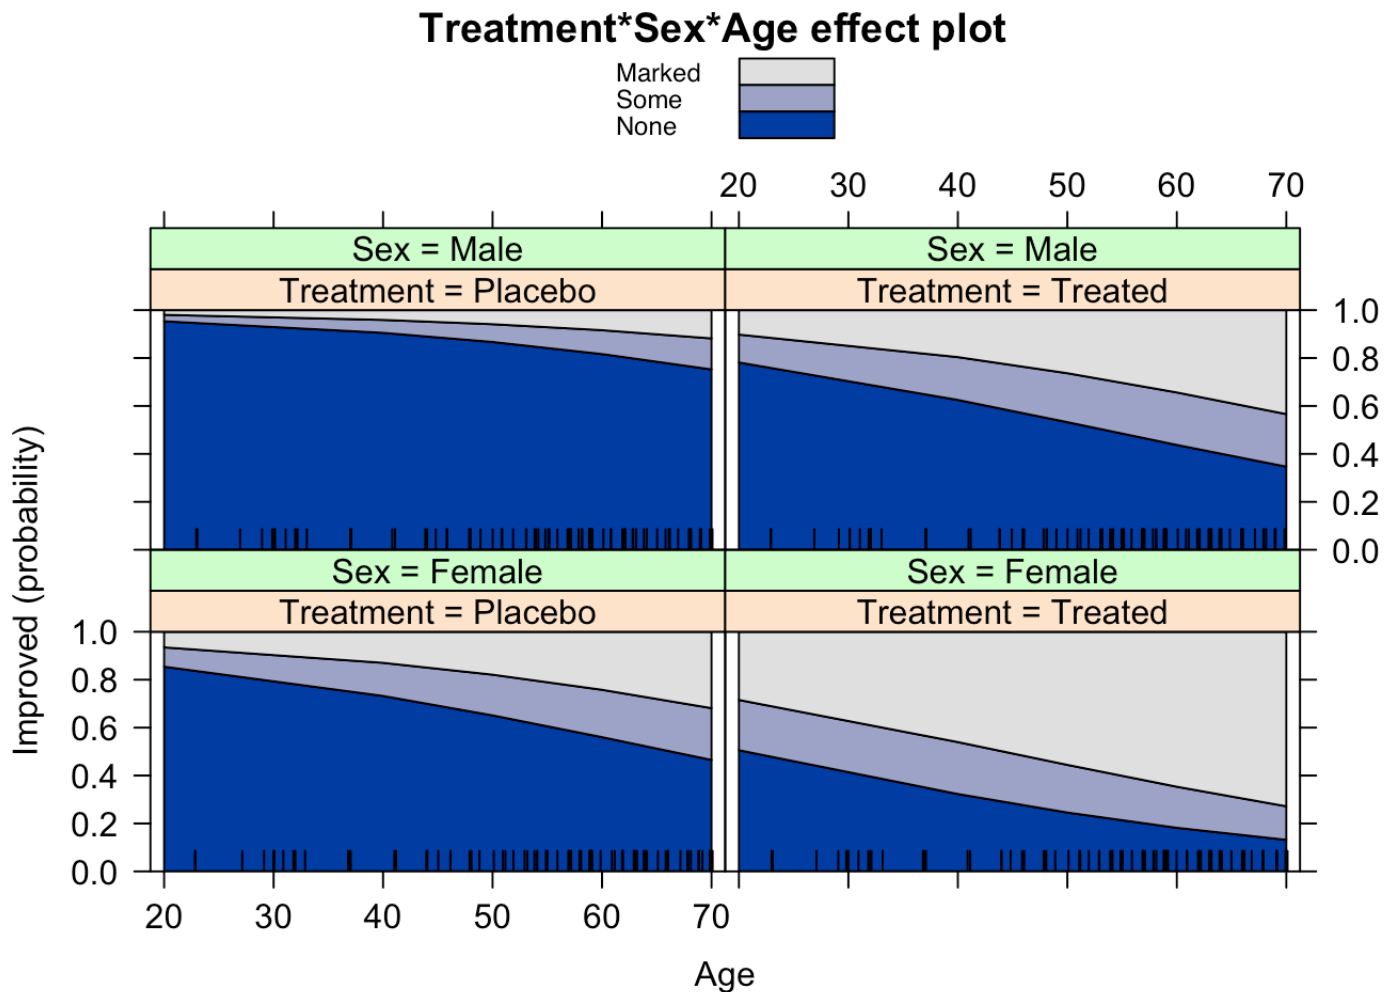


```
plot(Effect("Age", arth.polr), style='stacked',key.args=list(x=.55, y=.9))
```

## Age effect plot



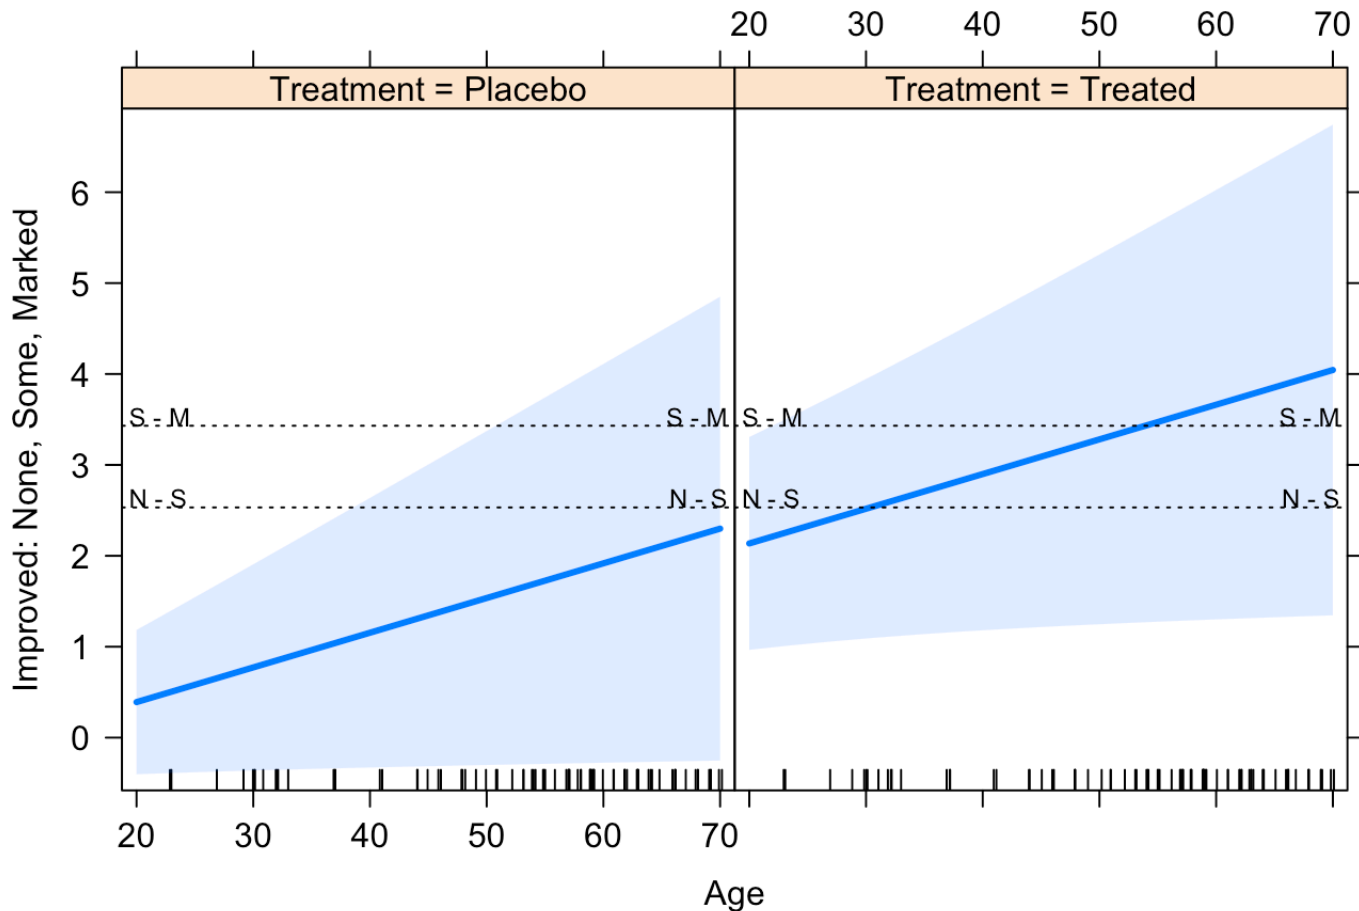
```
plot(Effect(c("Treatment", "Sex", "Age"), arth.polr), style="stacked", key.arg=list(x=.8, y=.9))
```



## Latent variable effect plot for the effects of Treatment and Age in the Arthritis data

```
plot(Effect(c("Treatment", "Age"), arth.polr, latent=TRUE), lwd=3)
```

## Treatment\*Age effect plot



## EXAMPLE 8.1: Women's labor force participation

```
library(car)
rm(Womenlf)
data("Womenlf", package="car")
#head(Womenlf)
```

```
Womenlf <- Womenlf
```

In this example, it makes sense to consider a first dichotomy (working) between women who are not working, vs. those who are (full time or part time). A second dichotomy (fulltime) contrasts full time work vs. part time work, among those women who are working at least part time. These two binary variables are created in the data frame using the `recode()` function from the `car` package

```
#Create dichotomies// using ifelse. Recode is another option which was not working in this case  
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:Hmisc':  
##  
##     src, summarize
```

```
## The following object is masked from 'package:car':  
##  
##     recode
```

```
## The following object is masked from 'package:MASS':  
##  
##     select
```

```
## The following objects are masked from 'package:stats':  
##  
##     filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
##     intersect, setdiff, setequal, union
```

```
working = ifelse(Womenlf$partic=='not.work',"no","yes")  
fulltime = ifelse(Womenlf$partic=='fulltime',"yes",  
                  ifelse(Womenlf$partic=='parttime',"no",NA)  
                  )  
Womenlf = cbind(Womenlf,working,fulltime)  
#Womenlf = cbind(Womenlf,fulltime)  
head(Womenlf)
```

```
##      partic hincome children  region working fulltime
## 1 not.work      15  present Ontario      no      <NA>
## 2 not.work      13  present Ontario      no      <NA>
## 3 not.work      45  present Ontario      no      <NA>
## 4 not.work      23  present Ontario      no      <NA>
## 5 not.work      19  present Ontario      no      <NA>
## 6 not.work       7  present Ontario      no      <NA>
```

```
with(Womenlf, table(partic, working))
```

```
##           working
## partic      no yes
## fulltime    0  66
## not.work  155   0
## parttime    0  42
```

```
with(Womenlf, table(partic, fulltime, useNA="ifany"))
```

```
##           fulltime
## partic      no yes <NA>
## fulltime    0  66   0
## not.work    0   0  155
## parttime   42   0   0
```

We proceed to fit two separate binary logistic regression models for the derived dichotomous variables. For the working dichotomy, we get the following results:

```
mod.working <- glm(working ~ hincome + children, family=binomial,data=Womenlf)
summary(mod.working)
```



```
##
## Call:
## glm(formula = working ~ hincome + children, family = binomial,
##      data = Womenlf)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6767  -0.8652  -0.7768   0.9292   1.9970
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    1.33583    0.38376   3.481  0.0005 ***
## hincome        -0.04231    0.01978  -2.139  0.0324 *
## childrenpresent -1.57565    0.29226  -5.391  7e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 356.15  on 262  degrees of freedom
## Residual deviance: 319.73  on 260  degrees of freedom
## AIC: 325.73
##
## Number of Fisher Scoring iterations: 4
```

```
mod.fulltime <- glm(fulltime ~ hincome + children, family=binomial,data=Womenlf)
summary(mod.fulltime)
```

```
##
## Call:
## glm(formula = fulltime ~ hincome + children, family = binomial,
##      data = Womenlf)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4047  -0.8678   0.3949   0.6213   1.7641
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    3.47777    0.76711   4.534 5.80e-06 ***
## hincome        -0.10727    0.03915  -2.740  0.00615 **
## childrenpresent -2.65146    0.54108  -4.900 9.57e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 144.34  on 107  degrees of freedom
## Residual deviance: 104.49  on 105  degrees of freedom
## (155 observations deleted due to missingness)
## AIC: 110.49
##
## Number of Fisher Scoring iterations: 5
```

Although these were fit separately, we can view this as a combined model for the three-level response, with the following coefficients

```
cbind(working=coef(mod.working), fulltime=coef(mod.fulltime))
```

```
##              working  fulltime
## (Intercept)    1.33582979  3.4777735
## hincome        -0.04230843 -0.1072679
## childrenpresent -1.57564843 -2.6514557
```

For both dichotomies, increasing income of the husband and the presence of young children decrease the log odds of a greater level of work. However, for those women who are working the effects of husband's income and children are greater on the choice between full time and part time work than they are for all women on the choice between working and not workin

```
LRtest <- function(model) c(LRchisq=(model$null.deviance - model$deviance),df=(model$
df.null - model$df.residual))
tab <- rbind(working=LRtest(mod.working),fulltime=LRtest(mod.fulltime))
tab <- rbind(tab, All = colSums(tab))
tab <- cbind(tab, pvalue = 1- pchisq(tab[,1], tab[,2]))
tab
```

```
##           LRchisq df      pvalue
## working  36.41835  2 1.235536e-08
## fulltime 39.84682  2 2.225215e-09
## All      76.26518  4 1.110223e-15
```

```
Anova(mod.working)
```

```
## Analysis of Deviance Table (Type II tests)
##
## Response: working
##           LR Chisq Df Pr(>Chisq)
## hincome      4.8264  1   0.02803 *
## children    31.3229  1  2.185e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Anova(mod.fulltime)
```

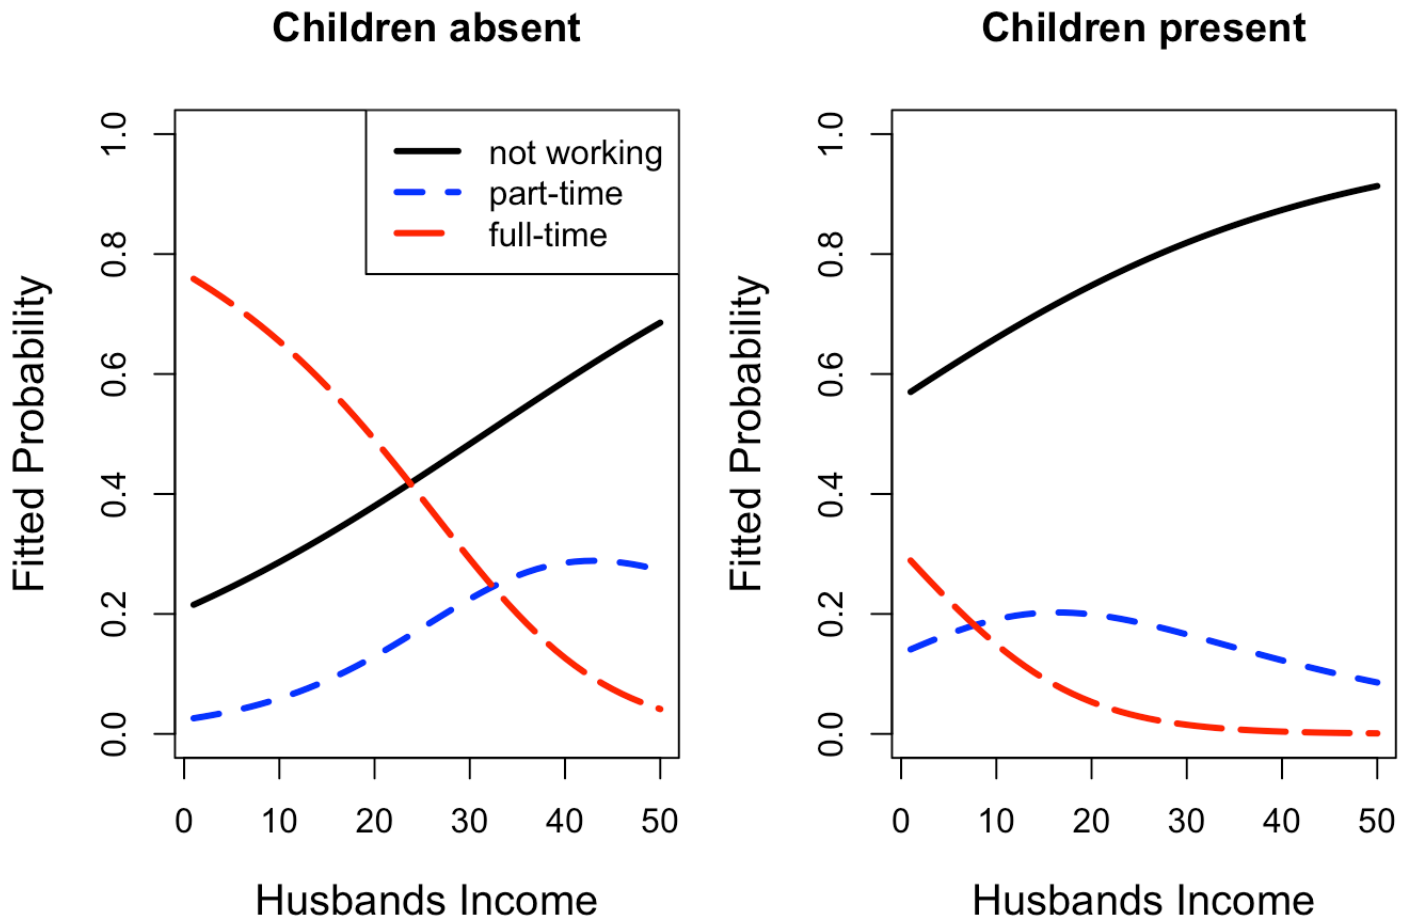
```
## Analysis of Deviance Table (Type II tests)
##
## Response: fulltime
##           LR Chisq Df Pr(>Chisq)
## hincome      8.981  1   0.002728 **
## children    32.136  1  1.437e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
predictors <- expand.grid(hincome=1:50,children=c('absent', 'present'))
fit <- data.frame(predictors,
                  p.working = predict(mod.working, predictors, type='response'),
                  p.fulltime = predict(mod.fulltime, predictors, type='response'),
                  l.working = predict(mod.working, predictors, type='link'),
                  l.fulltime = predict(mod.fulltime, predictors, type='link')
)
print(some(fit, 5), digits=3)
```

```
##      hincome children p.working p.fulltime l.working l.fulltime
## 18      18   absent    0.640    0.8245    0.574    1.547
## 29      29   absent    0.527    0.5907    0.109    0.367
## 34      34   absent    0.474    0.4578   -0.103   -0.169
## 37      37   absent    0.443    0.3796   -0.230   -0.491
## 96      46  present    0.101    0.0162   -2.186   -4.108
```

```
fit <- within(fit, { full <- p.working * p.fulltime
part <- p.working * (1 - p.fulltime)
not  <- 1 - p.working
})
```

```
op <- par(mfrow=c(1,2), mar=c(5,4,4,1)+.1)
Hinc <- 1:max(fit$hincome)
for ( kids in c("absent", "present") )
{
  dat <- subset(fit,children==kids)
  plot(range(Hinc),c(0,1),type="n",cex.lab=1.25,
        xlab="Husbands Income",ylab="Fitted Probability",
        main=paste("Children",kids))
  lines(Hinc, dat$not, lwd=3, col="black", lty=1)
  lines(Hinc, dat$part, lwd=3, col="blue", lty=2)
  lines(Hinc, dat$full, lwd=3, col="red", lty=5)
  if(kids=="absent")
  {
    legend("topright", lty=c(1,2,5), lwd=3, col=c("black", "blue", "red"),
          legend=c('not working', 'part-time', 'full-time'))
  }
}
```




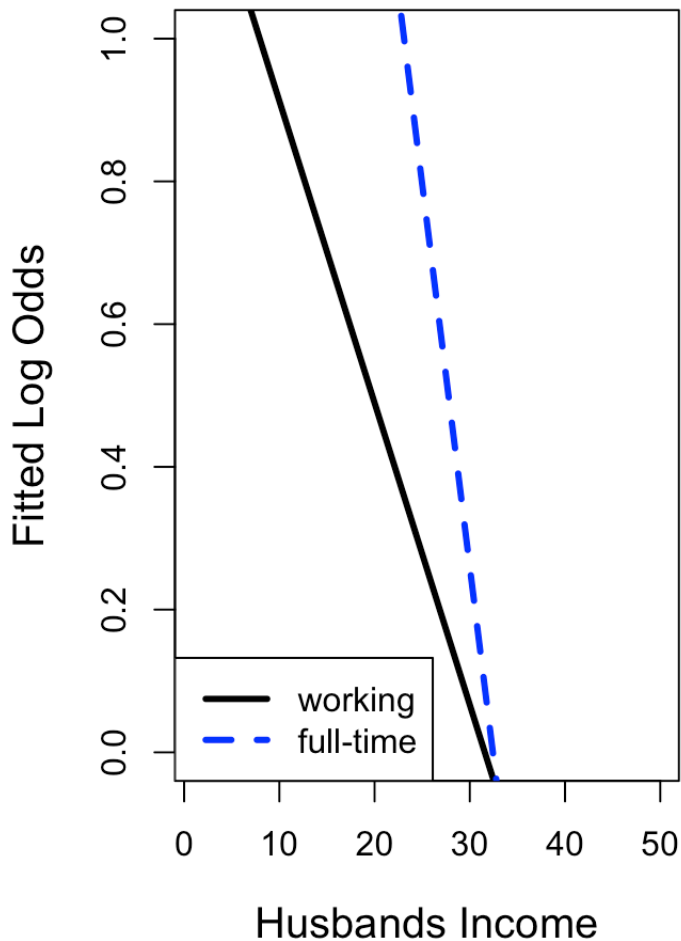
```
par(op)
```

We can see how that the decision not to work outside the home increases strongly with husband's income, and is higher when there are children present. As well, among working women, the decision to work full time as opposed to part time decreases strongly with husband's income, and is less likely with young children. Similarly, we plot the fitted logits for the two dichotomies in `l.working` and `l.fulltime` as shown below:

```

op <- par(mfrow=c(1,2), mar=c(5,4,1,1)+.1)
for ( kids in c("absent", "present") )
{
  dat <- subset(fit,children==kids)
  plot(range(Hinc),c(0,1),type="n",cex.lab=1.25,
        xlab="Husbands Income",ylab="Fitted Log Odds")
  lines(Hinc, dat$l.working, lwd=3, col="black", lty=1)
  lines(Hinc, dat$l.fulltime, lwd=3, col="blue", lty=2)
  text(25, 4.5, paste("Children", kids), cex=1.4)
  if (kids=="absent")
  {
    legend("bottomleft", lty=1:2, lwd=3, col=c("black", "blue"),
          legend=c('working', 'full-time'))
  }
}

```



These Graphs didn't come out as expected

# Generalized Logit Model

Response : Women working full time, part time or not working Variables: Husbands Income (Continuous), Children Present vs. Absent (Categorical)

```
levels(Womenlf$partic)
```

```
## [1] "fulltime" "not.work" "parttime"
```

```
# choose not working as baseline category
Womenlf$partic <- relevel(Womenlf$partic, ref="not.work")
```

We fit the main effects model for husband's income and children as follows. As we did with `polr()` (Section 8.1), specifying `Hess=TRUE` saves the Hessian and facilitates calculation of standard errors and hypothesis tests

```
library(nnet)
wlf.multinom <- multinom(partic ~ hincome + children,data=Womenlf, Hess=TRUE)
```

```
## # weights:  12 (6 variable)
## initial  value 288.935032
## iter   10 value 211.454772
## final   value 211.440963
## converged
```

The `summary()` method for “multinom” objects doesn't calculate test statistics for the estimated coefficients by default. The option `Wald=TRUE` produces Wald z-test statistics, calculated as  $z = \beta/SE(\beta)$

```
summary(wlf.multinom, Wald=TRUE)
```

```
## Call:
## multinom(formula = partic ~ hincome + children, data = Womenlf,
##          Hess = TRUE)
##
## Coefficients:
##          (Intercept)      hincome childrenpresent
## fulltime    1.982842 -0.097232073      -2.55860537
## parttime   -1.432321  0.006893838       0.02145558
##
## Std. Errors:
##          (Intercept)      hincome childrenpresent
## fulltime    0.4841789 0.02809599       0.3621999
## parttime    0.5924627 0.02345484       0.4690352
##
## Value/SE (Wald statistics):
##          (Intercept)      hincome childrenpresent
## fulltime    4.095266 -3.4607098      -7.06407045
## parttime   -2.417573  0.2939197       0.04574407
##
## Residual Deviance: 422.8819
## AIC: 434.8819
```

Notice that the coefficients, their standard errors and the Wald test z values are printed in separate tables. The first line in each table pertains to the logit comparing full time work with the not working reference level; the second line compares part time work against not working

For those who like p-values for significance tests, you can calculate these from the results returned by the `summary()` method in the `Wald.ratios` component, using the standard normal asymptotic approximation

```
stats <- summary(wlf.multinom, Wald=TRUE)
z <- stats$Wald.ratios
p <- 2 * (1 - pnorm(abs(z)))
zapsmall(p)
```

```
##          (Intercept)      hincome childrenpresent
## fulltime    0.0000422 0.0005388       0.0000000
## parttime    0.0156244 0.7688193       0.9635142
```

```
wlf.multinom2 <- multinom(partic ~ hincome * children, data=Womenlf, Hess=TRUE)
```



```
## # weights: 15 (8 variable)
## initial value 288.935032
## iter 10 value 210.797079
## final value 210.714841
## converged
```

```
Anova(wlf.multinom2)
```

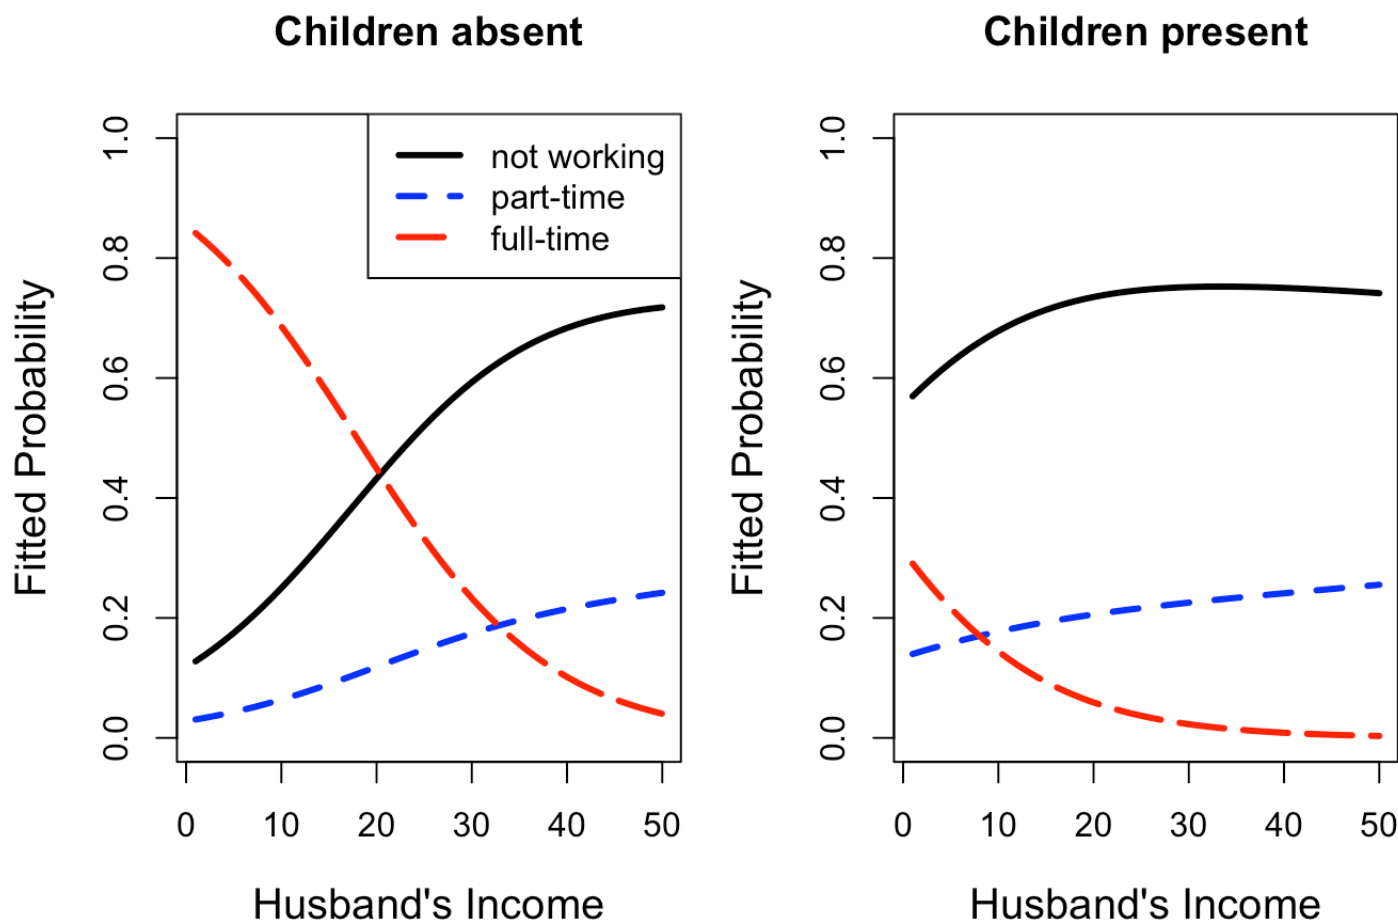
```
## Analysis of Deviance Table (Type II tests)
##
## Response: partic
##
##          LR Chisq Df Pr(>Chisq)
## hincome      15.153  2  0.0005123 ***
## children      63.559  2  1.579e-14 ***
## hincome:children  1.452  2  0.4837815
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Full model plots of the fitted values can be plotted as shown earlier in Example 8.1: obtain the fitted values over a grid of the predictors and plot these

```
predictors <- expand.grid(hincome=1:50,
                          children=c('absent', 'present'))
fit <- data.frame(predictors, predict(wlf.multinom, predictors, type='probs'))
```

Plotting these fitted values gives the plot shown in Figure 8.12.

```
op <- par(mfrow=c(1,2), mar=c(5,4,4,1)+.1)
Hinc <- 1:max(fit$hincome)
for ( kids in c("absent", "present") ) {
  dat <- subset(fit, children==kids)
  plot( range(Hinc), c(0,1), type="n", cex.lab=1.25,
        xlab="Husband's Income", ylab='Fitted Probability',
        main = paste("Children", kids))
  lines(Hinc, dat$not.work, lwd=3, col="black", lty=1)
  lines(Hinc, dat$parttime, lwd=3, col="blue", lty=2)
  lines(Hinc, dat$fulltime, lwd=3, col="red", lty=5)
  if (kids=="absent") {
    legend("topright", lty=c(1,2,5), lwd=3, col=c("black", "blue", "red"), legend=c('not working', 'part-time', 'full-time'))
  }
}
```



The effects package has special methods for “multinom” models. It treats the response levels in the order given by levels(), so before plotting we use ordered() to arrange levels in their natural order. The update() method provides a simple way to get a new fitted model; in the call, the model formula. ~  
 .meanstofitthesamemodelasbefore,i.e.,partic ~ hincome + children

```
Womenlf$partic <- ordered(Womenlf$partic,
                          levels=c('not.work', 'parttime', 'fulltime'))
wlf.multinom <- update(wlf.multinom, . ~ .)
```

```
## # weights: 12 (6 variable)
## initial value 288.935032
## iter 10 value 211.454772
## final value 211.440963
## converged
```

As illustrated earlier, you can use plot(allEffects(model), ...) to plot all the high-order terms in the model, either with separate curves for each response level (style=“lines”) or as cumulative filled polygons (style=“stacked”). Here, we simply plot the effects for the combinations of husband’s income and children in stacked style, giving a

plot (Figure 8.13) that is analogous to the full-model plot shown in Figure 8.12.

```
plot(Effect(c("hincome", "children"), wlf.multinom),
     style="stacked", key.args=list(x=.05, y=.9))
```

### hincome\*children effect plot

